

Proactive Support through floating window in Information Foraging with User Attention Modeling via Eye-tracking

Sumin Hong
shong6@nd.edu
University of Notre Dame
Notre Dame, Indiana, USA

Brandon Forseth
bforseth@nd.edu
University of Notre Dame
Notre Dame, Indiana, USA

Toby Jia-Jun Li
toby.j.li@nd.edu
University of Notre Dame
Notre Dame, Indiana, USA




Figure 1: A proactive support system leveraging user attention modeling via eye-tracking in information foraging tasks. (a) The system accumulates user intention across different search and browsing contexts. (b) Based on attention states, the system proactively suggests relevant content. (c) Users can trigger suggestions through interactions. (d) The system automatically highlights previously attended content to reduce cognitive overload. The floating support windows adaptively respond to inferred user goals, aiding efficient navigation in complex information spaces.

Abstract

In complex online information foraging tasks, users often face cognitive overload due to fragmented interfaces and the limits of working memory. While large language models (LLMs) offer content-based assistance, they often miss users' true information needs by relying solely on browsing history. We introduce an attention-driven system that models user gaze as a grounding signal to infer cognitive states and evolving mental context. Leveraging real-time eye-tracking, our system delivers proactive interventions—contextual zoom-in, information pruning, and distraction recovery—through a lightweight floating UI. Designed to be non-intrusive and seamlessly integrated into users' browsing workflow, the interface enhances understanding without disrupting task flow or requiring explicit prompts. By accumulating attention over time, the system builds a dynamic knowledge graph to support multi-session exploration. This position paper shows how combining attention modeling with

proactive support design can reduce cognitive load, support comprehension, and enable more intuitive, context-sensitive information foraging.

CCS Concepts

• **Human-centered computing;**

Keywords

Human-AI Collaboration, User Interface, Sensemaking, Generative AI, Large Language Models, Eye-tracking, Multi-Modality, Cognitive Modeling

ACM Reference Format:

Sumin Hong, Brandon Forseth, and Toby Jia-Jun Li. 2025. Proactive Support through floating window in Information Foraging with User Attention Modeling via Eye-tracking. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (CHI 25 UI Workshop)*. ACM, New York, NY, USA, 7 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

In an online environment, users frequently engage in information tasks that involve exploring, understanding, and synthesizing unfamiliar domains or concepts to develop a comprehensive understanding or make informed decisions. This iterative process exposes users to a large volume of information, primarily accessed through web browsing, which requires repeatedly navigating between multiple web pages and windows. Figure 1 illustrates this scenario, showing how users visit various websites while searching for an

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

CHI 25 UI Workshop, Yokohama, Japan

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-XXXX-X/2018/06

<https://doi.org/XXXXXXX.XXXXXXX>

audio resolution enhancing API with a specific purpose. A key challenge in this process is the limited capacity of working memory when processing unfamiliar information [36, 48], compounded by extraneous cognitive load [27] introduced by disjointed web interfaces that fragment information across stacked windows and interrupt users' non-linear foraging patterns.

To address these challenges, prior research has proposed systems that proactively re-organize or summarize content [12, 22, 32?, 33], often layered on top of the current web interface. Other studies support information foraging in multi-criteria decision-making [10, 31, 32] or externalize users' mental contexts [11, 29] using similar methods. Many of these systems increasingly incorporate large language models (LLMs) to support content understanding and retrieval.

However, simply tuning LLMs based on users' browsing history is often insufficient. Users frequently engage in "selective reading" [54, 55], focusing only on parts of a page they find relevant while ignoring the rest. This creates a disconnect: what users visit isn't always what they read, and what they read isn't always what they understand. As a result, LLMs that rely solely on browsing logs may misidentify users' true information needs, resulting in either irrelevant suggestions [46] or misleading over-accommodations [7, 43–45].

To bridge this gap, we propose modeling user attention as an act of grounding [13]. Understanding user intention and status can make LLMs serve more contextual and helpful support. For effective attention modeling, we utilize eye-tracking as a foundation. Gaze is a key nonverbal signal for intent [3, 17, 25], and prior work has shown its utility in understanding user behaviors [14] and cognitive states [21, 30, 51, 52]. By processing gaze data into attention states, our system can deliver more context-sensitive interventions.

Building on this modeling method, we introduce an integrated system that combines attention-driven proactive support with knowledge graph construction. As shown in Figure 1(a), our system accumulates attention history over time to support context-aware interventions across sessions. The floating window interface delivers timely suggestions—such as contextual explanations or content highlights—without disrupting workflow or requiring prompt crafting. Unlike previous methods, our approach filters and structures information using real-time attention data rather than static browsing history, enabling more relevant and cognitively aligned support.

We implement three proactive interventions based on user attention states: (1) *Contextual Zoom-in*, triggered by sustained fixations, provides compact explanations via floating cards; (2) *Information Pruning*, activated by dispersed gaze patterns, highlights content relevant to prior attention; and (3) *Distraction Recovery*, which helps users regain focus after detecting repetitive or aimless navigation.

In summary, this paper explores how attention modeling via eye tracking can enhance web interfaces for information foraging. Our system combines proactive support through floating UI elements with a dynamic knowledge graph grounded in attention data. By doing so, we address the limitations of working memory and extraneous cognitive load while improving long-term information retrieval and understanding in multi-session exploration.

2 Related Works

2.1 Sensemaking in Online Information Foraging

Existing research on supporting sensemaking in online information foraging has addressed these challenges from various perspectives. Some studies [4, 18, 38, 42] are focused on information collection step, providing an overview of the information space not to be confused by encountering unfamiliar things. For the organizing information step, externalizing their mental contexts, such as intention or interests, instantaneously was effective without any methods out of the web browsers to arrange their understanding, which occurs in additional steps [11, 29]. Research addressing the online decision-making process [10, 31, 32] handled the issue from multiple criteria and options existing. Other studies [12, 15, 22, 32, 33] proposed proactive approaches to structure the information by tracking the contents of interests. Most of these tasks have approached representation by providing re-organized visual or textual support for side views of the web interface.

2.2 Issues of Limited Working Memory Capacity and Knowledge Retrieving

When exploring unfamiliar concepts, participants frequently encounter the rabbit hole phenomenon [19, 35], engrossed in chains of nested concepts diverting them from their initial research objectives. This phenomenon is particularly prevalent in the early stages of sensemaking, as users attempt to navigate, filter, and comprehend complex topics through reading and understanding information from different resources. Additionally, this problem is compounded by context-switching between different resources or interruptions which easily happen in web searching [10]. This goal displacement disrupts the transition from information foraging to meaningful schema construction, as users become trapped in expanding information collection rather than progressing to the synthesis and organization phase of sensemaking. The challenge is further exacerbated by physical breaks between tasks, which blur previously acquired knowledge and goals, requiring additional cognitive resources to reestablish context upon returning to the task. These issues increase cognitive workload due to our limited working memory capacity. To retrieve and reuse knowledge forgotten due to this cognitive load, prior research has proposed reconstructing past contexts through *histories*, such as tracking desktop tasks [16], summarizing browsing activity [15, 23], and capturing screenshots [24, 40] to reconstruct the context of past computer work. However, these approaches primarily provide simple visual or textual histories. Additionally, personal knowledge graphs with question answering can be one way to retrieve their previous knowledge [9].

2.3 Relation between UI, Cognitive States, and Eye-tracking

Eye-tracking is a method that measures where, how long, and in what manner a person looks at a given stimulus using an eye tracker. It serves as a tangible way of representing attention. Various eye-tracking metrics [5], such as fixation, saccades, blinks, and pupillometry, provide insights into different aspects of gaze behavior. Research leveraging these eye-tracking metrics has been

conducted across multiple domains to examine selection, focus, and cognitive states. Eye-tracking has been widely utilized in cognitive-demanding fields such as software engineering [39], machine operation [50], decision making [8] and driving [56] to understand user behavior patterns and assess cognitive load. Based on these capabilities, several studies have attempted to predict users' cognitive states [6, 20, 53], while others have focused on generating scanpaths for given UI designs [26]. Data collection remains a crucial component of these modeling studies, with ongoing research exploring how to collect reliable data using simple equipment [37]. Given these characteristics, researchers have leveraged eye-tracking methods to inform user interface design and improvements [14], creating a valuable feedback loop between cognitive understanding and interaction design.

3 Attention Modeling

Gaze movement is a very natural behavior; it has been used to measure cognitive load, how users' attention goes, etc. Given this characteristic, our research team explores the way of modeling attention states through eye-tracking data, which lowers the burden that comes from trying additional behaviors such as moving to another window, expressing their thoughts as text, or being used to a transformed interface. As a goal of our system is to conserve the natural searching process as much as possible, we draft a modeling method for user attention through eye-tracking techniques, which can be fundamental to the system.

In this section, we are talking about data collection methods, creating a taxonomy, and modeling methods.

3.1 Data Collection

We have developed a protocol to collect eye movement data from participants. The protocol includes six carefully designed tasks that require multiple searching and thinking processes. Each participant is asked to answer questions related to these tasks. For each task, we provide participants with approximately 10 pre-selected websites in the form of a hyperlink list. These websites contain multiple text paragraphs with relevant information from sources such as articles and online forums.

Participants are randomly assigned two questions to answer through web searching. After task completion, we process the collected data and screen recordings into short clips of 8-10 seconds each, enhanced with scanpath visualization. This visualization helps participants recall their cognitive states during the task. We ask participants to watch these clips and assign a single attention state label to each clip based on their memory of the experience. While we provide our initial taxonomy to guide the labeling process, participants can also suggest new labels if they feel none of the existing categories adequately represent their attentional state. Further details regarding the taxonomy creation process are in Section 3.2.

For eye movement tracking, we use a monitor-mounted Tobii Fusion Pro eye tracker. Since this equipment requires a stable environment to capture gaze accurately, we conduct all data collection in a controlled lab-setting with consistent lighting. Participants are seated in non-swivel chairs and instructed to maintain a stable posture throughout the session to ensure data quality.




Figure 2: Initial taxonomy of attention states. It will be given to participants to facilitate the labeling process, but should be changed through data collecting process.

3.2 Creating Taxonomy

For classification tasks, a proper taxonomy is essential. For our modeling approach, we propose an initial taxonomy of eight distinct states based on how users explore web pages: two initial navigation states and six subsequent exploration states (Figure 2). While this taxonomy cannot yet fully represent users' attentional states, it will be reconstructed through a data collection study where participants directly label their own states.

To label a particular eye-tracking segment, we ask participants to consider the sequential structure of attention states shown in Figure 2. The sequential structure first distinguishes segments by whether the user is in an initial navigation phase or in a deeper exploration phase. For example, in a segment where a participant was initially scanning a page and then engaged in deep reading of specific content, the appropriate label would be "Wandering" as this represents the primary attentional state during that segment. In cases where two states are appropriate and fall under the same stage (e.g., if the participant was both reading carefully and thinking specific contents), they are asked to pick the state in which they spent the majority of the time. These issues arise because we collect a single attention state label for each eye-tracking segment.

After entering a webpage through a search engine, users typically exhibit one of two initial states depending on whether they have a clear navigational purpose: either general top-left exploration or direct navigation to specific information. Subsequently, as users explore webpage content, we categorize their behavior into six states based on comprehension level and degree of focus. These categories can be analyzed using several eye-tracking metrics outlined in Table 1.

3.3 Eye Movement Data and Modeling

The collected eye movement dataset will comprise numerous labeled segments, with each segment containing approximately 1024 data points. Instead of directly using raw x and y coordinates, we convert these data points into low-level gaze events (detailed in Table 1) that can be meaningfully linked to users' attentional states. This conversion strategy helps mitigate the challenges posed by varying UI layouts across multiple websites.

Metric	Description
Fixation [41]	Highlights sections that users selectively focus on, indicating concentrated reading behavior.
Saccade [41]	Measures gaze transition speed; rapid and repeated scanning over a broad area may indicate cognitive overload or difficulty in information processing.
Pupillometry [49]	Quantifies attention by measuring pupil dilation per unit time to identify the not-thinking state.
Blink Rate [34]	Tracks the number of blinks per unit time; a decreased blink rate is often associated with increased cognitive load and attention.

Table 1: Primary low-level gaze events and their application to attention modeling.

Table 1 describes the low-level metrics that provide insights into user engagement and facilitate the identification of key focus areas in gaze scan paths. Specifically, fixation and saccade metrics enable us to infer reading patterns that indicate deeper cognitive processing. Complementing these behavioral measures, psychophysiological metrics such as pupil dilation and blink rate offer additional physiological indicators of cognitive load and attentional shifts.

Within each webpage, we designate UI elements such as navigation bars, search fields, and interactive buttons as areas of interest (AOIs). These AOIs facilitate the measurement of gaze distribution patterns, which directly contribute to the weighting of knowledge base elements. We also incorporate the total time spent on each webpage to determine its overall significance. These computed weightings play a crucial role in constructing the personalized knowledge base described in Section 4.2.

This approach enables us to refine the importance of AOIs based on sustained focus patterns rather than isolated gaze instances. Additionally, we derive higher-level features including fixation duration distributions, saccade amplitudes, scanpath complexity metrics, and transition probabilities between AOIs to create a comprehensive feature set for our attention classification model.

4 System Design

4.1 Proactive Support based on User Attention

The automatic support powered by attention modeling in Figure 1(a) represents a significant advancement over previous practices, requiring less effort and mental workload to acquire LLMs support. For our main interface, we propose a floating window design to address key challenges in conventional LLMs interactions. Traditional interactions that involve switching windows and manually crafting prompts often disrupt the natural information-seeking flow, increasing cognitive load whereas the purpose is to reduce it. Our floating window interface minimizes context switching costs while maintaining design versatility. Industry solutions such as Arc [1] and Sider [2] have adopted similar floating interfaces to preserve

information-seeking flow. Our research extends this paradigm by implementing a proactive mechanism driven by user attention patterns. The system continuously monitors attention states during information foraging and triggers predefined support interventions at the moments. While this approach shows promise, the user perception of such interventions remains largely unexplored.

The effectiveness of proactive assistance depends critically on timing—providing support precisely when users need it most while avoiding unnecessary interruptions. To identify the optimal balance between assistance and distraction, we propose three proactive mechanisms tailored to specific attention states:

Contextual Zoom-in: This support activates when users stare at specific screen areas, interpreted as needing additional information to develop their thoughts. For example, when a user is detected fixating on a technical term like "transformer architecture" for more than 3 seconds without scrolling or clicking (indicating cognitive engagement without comprehension), the system triggers a small floating card with a simplified explanation: "Transformers are neural networks that process sequences using attention mechanisms. Need more details?" This allows users to explore topics deeper through conversation without disrupting their reading flow.

Information Pruning: This intervention responds to detected wandering attention, often indicating high mental workload. Consider a scenario where eye-tracking identifies rapid scanning across multiple sections of a dense research paper. The system highlights the three most relevant paragraphs based on previous reading patterns and search query intent, temporarily reducing the visual prominence of peripheral content. This guidance redirects focus to pertinent information, reducing cognitive load while maintaining access to the complete document.

Distraction Recovery: This support activates when detecting repetitive viewing patterns in the same areas, suggesting lost concentration or confusion. When the system observes a user repeatedly reading the same paragraph about climate modeling three times in succession, it might gently prompt: "I notice you've reviewed this section multiple times. Are you looking for something specific about climate models? Your original search was about 'effects of carbon emissions on ocean temperatures'." This contextual reminder helps users recover their original search purpose after losing focus. All support mechanisms follow a two-step approach: the system automatically suggests concise information through a floating window, and users can select whether they wish to explore further. We are planning to refine these mechanisms through iterative modeling and user feedback.

However, two significant challenges are expected. Firstly, false positive errors in attention state classification may lead to inappropriate interventions that users perceive as disruptive. We are exploring rule-based heuristics applied on top of classification results to mitigate this risk. Secondly, the displayed content may not always align with users' actual information needs. To address this challenge, we are developing a knowledge graph approach as a closer proxy for user intention and corresponding knowledge requirements.




Figure 3: A knowledge graph synthesized from users' search history. It can provide an organized, linear browsing history that reflects the individual's search intent and focus at the time.

4.2 Synthesizing Knowledge Graph of Browsing History

To enable effective proactive support as described in Section 4.1, building a knowledge graph by filtering and accumulating information from browsing history can be beneficial (Figure 1(a), (d)). Moving beyond the limitations of previous less-grounded LLMs, integrating attention modeling with this graph-based knowledge base can provide more effective information retrieval. This knowledge graph can be constructed by prioritizing information based on user attention patterns. Concepts that users find confusing or significant can be emphasized through highlighting, while less relevant information can be deprioritized.

Moreover, the knowledge graph itself can be a well-organized way of representing previous browsing history while capturing user intention. Figure 3 depicted the graph. Imagine resuming a search for an audio enhancing API that you started three days or three hours ago. We typically forget the context and methods of our previous searches. Offering the relevant browsing history in visual graph form representing the relationships between explored topics can aid memory retrieval, rather than linearly listed visited website history.

However, a potential drawback of prioritizing information stems from the tendency of domain novices to oversimplify problems by selecting only a subset of features to address, causing them to miss important problem features or relationships [28, 47]. This situation could lead to knowledge graphs that are misaligned with what users actually need to acquire through the browsing process. Our research team aims to address this challenge through proper weighting between attention patterns and domain knowledge, ensuring that personalization enhances rather than limits information comprehension.

5 Conclusion

This position paper proposes an attention-driven approach to enhance sensemaking in complex online environments. By modeling user attention through eye-tracking, our system captures real-time cognitive signals that reflect users' evolving focus, intent, and comprehension. This fine-grained attention modeling serves as the foundation for context-aware support that goes beyond traditional prompt-based interactions. Building on this foundation, our system delivers proactive interventions—such as contextual explanations, content highlighting, and distraction recovery—through a lightweight floating UI that minimizes workflow disruption. These interventions are driven not by browsing history alone, but by dynamic attention patterns, allowing for more precise and cognitively aligned assistance. Furthermore, our system constructs a knowledge graph grounded in users' attention histories, enabling cross-session continuity and richer context retention. This multi-modal approach, centered on user cognition, lays the groundwork for more intuitive, responsive, and human-centered generative AI systems. Looking ahead, we plan to incorporate additional modalities such as mouse movement and biometrics, and to explore more natural interaction methods for grounding AI in users' lived workflows.

References

- [1] 2025. Arc from The Browser Company. <https://arc.net>
- [2] 2025. Sider: ChatGPT Sidebar + GPT-4.1, Claude 3.5, Gemini 2.0 & AI Tools. <https://sider.ai/en>
- [3] Anna Belardinelli. 2024. Gaze-Based Intention Estimation: Principles, Methodologies, and Applications in HRI. *J. Hum.-Robot Interact.* 13, 3 (Sept. 2024), 31:1–31:30. <https://doi.org/10.1145/3656376>
- [4] Krishna Bharat. 2000. SearchPad: explicit capture of search context to support Web search. *Comput. Netw.* 33, 1 (June 2000), 493–501. [https://doi.org/10.1016/S1389-1286\(00\)00047-5](https://doi.org/10.1016/S1389-1286(00)00047-5)
- [5] M. Borys and M. Plechawska-Wójcik. 2017. Eye-tracking metrics in perception and visual attention research. <https://www.semanticscholar.org/paper/Eye-tracking-metrics-in-perception-and-visual-Borys-Plechawska-W%3B3jck/bee219d15c336341d87c4be3d1cfcc327b0cefd>
- [6] Duncan P. Brumby and Andrew Howes. 2004. Good Enough But I'll Just Check: Web-page Search as Attentional Refocusing. <https://www.semanticscholar.org/paper/Good-Enough-But-I'll-Just-Check%3A-Web-page-Search-as-Brumby-Howes/feb8c3f76eefe9c588887bba5f40bafc17f8a8b0>
- [7] Alice Cai, Ian Arawjo, and Elena L. Glassman. 2024. Antagonistic AI. <https://doi.org/10.48550/arXiv.2402.07350> arXiv:2402.07350 [cs].
- [8] Nora Castner, Lubaina Arsiwala-Scheppach, Sarah Mertens, Joachim Krois, Enkeleda Thaqi, Enkeleda Kasneci, Siegfried Wahl, and Falk Schwendicke. 2024. Expert gaze as a usability indicator of medical AI decision support systems: a preliminary study. *npj Digital Medicine* 7, 1 (July 2024), 1–10. <https://doi.org/10.1038/s41746-024-01192-8> Publisher: Nature Publishing Group.
- [9] Prantika Chakraborty and Debarshi Kumar Sanyal. 2024. A Personal Knowledge Graph for Researchers. In *Proceedings of the 15th Annual Meeting of the Forum for Information Retrieval Evaluation (FIRE '23)*. Association for Computing Machinery, New York, NY, USA, 129–132. <https://doi.org/10.1145/3632754.3632773>
- [10] Joseph Chee Chang, Nathan Hahn, and Aniket Kittur. 2020. Mesh: Scaffolding Comparison Tables for Online Decision Making. In *Proceedings of the 33rd Annual ACM Symposium on User Interface Software and Technology*. ACM, Virtual Event USA, 391–405. <https://doi.org/10.1145/3379337.3415865>
- [11] Joseph Chee Chang, Nathan Hahn, Adam Perer, and Aniket Kittur. 2019. Search-Lens: composing and capturing complex user interests for exploratory search. In *Proceedings of the 24th International Conference on Intelligent User Interfaces (IUI '19)*. Association for Computing Machinery, New York, NY, USA, 498–509. <https://doi.org/10.1145/3301275.3302321>
- [12] Joseph Chee Chang, Yongsung Kim, Victor Miller, Michael Xieyang Liu, Brad A. Myers, and Aniket Kittur. 2021. Tabs.do: Task-Centric Browser Tab Management. In *The 34th Annual ACM Symposium on User Interface Software and Technology (UIST '21)*. Association for Computing Machinery, New York, NY, USA, 663–676. <https://doi.org/10.1145/3472749.3474777>
- [13] Herbert H. Clark and Susan E. Brennan. 1991. Grounding in communication. In *Perspectives on socially shared cognition*. American Psychological Association, Washington, DC, US, 127–149. <https://doi.org/10.1037/10096-006>

- [14] Edward Cutrell and Zhiwei Guan. 2007. What are you looking for? an eye-tracking study of information usage in web search. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '07)*. Association for Computing Machinery, New York, NY, USA, 407–416. <https://doi.org/10.1145/1240624.1240690>
- [15] Mira Dontcheva, Steven M. Drucker, Geraldine Wade, David Salesin, and Michael F. Cohen. 2006. Summarizing personal web browsing sessions. In *Proceedings of the 19th annual ACM symposium on User interface software and technology (UIST '06)*. Association for Computing Machinery, New York, NY, USA, 115–124. <https://doi.org/10.1145/1166253.1166273>
- [16] Anton N. Dragunov, Thomas G. Dietterich, Kevin Johnsrude, Matthew McLaughlin, Lida Li, and Jonathan L. Herlocker. 2005. TaskTracer: a desktop environment to support multi-tasking knowledge workers. In *Proceedings of the 10th international conference on Intelligent user interfaces (IUI '05)*. Association for Computing Machinery, New York, NY, USA, 75–82. <https://doi.org/10.1145/1040830.1040855>
- [17] N. J. Emery. 2000. The eyes have it: the neuroethology, function and evolution of social gaze. *Neuroscience & Biobehavioral Reviews* 24, 6 (Aug. 2000), 581–604. [https://doi.org/10.1016/S0149-7634\(00\)00025-7](https://doi.org/10.1016/S0149-7634(00)00025-7)
- [18] Raymond Fok, Nedim Lipka, Tong Sun, and Alexa F Siu. 2024. Marco: Supporting Business Document Workflows via Collection-Centric Information Foraging with Large Language Models. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems (CHI '24)*. Association for Computing Machinery, New York, NY, USA, 1–20. <https://doi.org/10.1145/3613904.3641969>
- [19] Amy Rae Fox, Robert Kaufman, and James D Hollan. [n. d.]. *Falling Down the Rabbit Hole: Exploratory Search and The Case for Personal Information Environments*. ([n. d.]).
- [20] Jarod P. Hartley. 2024. Task Classification using Eye Movements and Graph Neural Networks. In *Proceedings of the 2024 Symposium on Eye Tracking Research and Applications (ETRA '24)*. Association for Computing Machinery, New York, NY, USA, 1–3. <https://doi.org/10.1145/3649902.3655097>
- [21] John M. Henderson, Svetlana V. Shinkareva, Jing Wang, Steven G. Luke, and Jenn Olejarczyk. 2013. Predicting Cognitive State from Eye Movements. *PLOS ONE* 8, 5 (May 2013), e64937. <https://doi.org/10.1371/journal.pone.0064937> Publisher: Public Library of Science.
- [22] Andrew Hogue and David Karger. 2005. Thresher: automating the unwrapping of semantic content from the World Wide Web. In *Proceedings of the 14th international conference on World Wide Web (WWW '05)*. Association for Computing Machinery, New York, NY, USA, 86–95. <https://doi.org/10.1145/1060745.1060762>
- [23] Donghan Hu and Sang Won Lee. 2020. ScreenTrack: Using a Visual History of a Computer Screen to Retrieve Documents and Web Pages. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems (CHI '20)*. Association for Computing Machinery, New York, NY, USA, 1–13. <https://doi.org/10.1145/3313831.3376753>
- [24] Donghan Hu and Sang Won Lee. 2022. Scrapbook: Screenshot-Based Bookmarks for Effective Digital Resource Curation across Applications. In *Proceedings of the 35th Annual ACM Symposium on User Interface Software and Technology (UIST '22)*. Association for Computing Machinery, New York, NY, USA, 1–13. <https://doi.org/10.1145/3526113.3545678>
- [25] Chien-Ming Huang, Sean Andrisc, Allison Saupé, and Bilge Mutlu. 2015. Using gaze patterns to predict task intent in collaboration. *Frontiers in Psychology* 6 (July 2015). <https://doi.org/10.3389/fpsyg.2015.01049> Publisher: Frontiers.
- [26] Yue Jiang, Zixin Guo, Hamed Rezazadegan Tavakoli, Luis A. Leiva, and Antti Oulasvirta. 2024. EyeFormer: Predicting Personalized Scanpaths with Reasoner-Guided Reinforcement Learning. In *Proceedings of the 37th Annual ACM Symposium on User Interface Software and Technology (UIST '24)*. Association for Computing Machinery, New York, NY, USA, 1–15. <https://doi.org/10.1145/3654777.3676436>
- [27] Slava Kalyuga. 2011. Cognitive Load Theory: How Many Types of Load Does It Really Need? *Educational Psychology Review* 23, 1 (March 2011), 1–19. <https://doi.org/10.1007/s10648-010-9150-7>
- [28] Minchi C. Kim and Michael J. Hannafin. 2011. Scaffolding problem solving in technology-enhanced learning environments (TELEs): Bridging research and theory with practice. *Computers & Education* 56, 2 (Feb. 2011), 403–417. <https://doi.org/10.1016/j.compedu.2010.08.024>
- [29] Andrew Kuznetsov, Joseph Chee Chang, Nathan Hahn, Napol Rachatasumrit, Bradley Breneisen, Julina Coupland, and Aniket Kittur. 2022. Fuse: In-Situ Sensemaking Support in the Browser. In *Proceedings of the 35th Annual ACM Symposium on User Interface Software and Technology (UIST '22)*. Association for Computing Machinery, New York, NY, USA, 1–15. <https://doi.org/10.1145/3526113.3545693>
- [30] Tamera Lanham, Anna Chen, Ansh Radhakrishnan, Benoit Steiner, Carson Denison, Danny Hernandez, Dustin Li, Esin Durmus, Evan Hubinger, Jackson Kernion, Kamilé Lukošiušė, Karina Nguyen, Newton Cheng, Nicholas Joseph, Nicholas Schiefer, Oliver Rausch, Robin Larson, Sam McCandlish, Sandipan Kundu, Saurav Kadavath, Shannon Yang, Thomas Henighan, Timothy Maxwell, Timothy Telleen-Lawton, Tristan Hume, Zac Hatfield-Dodds, Jared Kaplan, Jan Brauner, Samuel R. Bowman, and Ethan Perez. 2023. Measuring Faithfulness in Chain-of-Thought Reasoning. <https://doi.org/10.48550/arXiv.2307.13702> arXiv:2307.13702 [cs].
- [31] Michael Xieyang Liu, Jane Hsieh, Nathan Hahn, Angelina Zhou, Emily Deng, Shaun Burley, Cynthia Taylor, Aniket Kittur, and Brad A. Myers. 2019. Unakite: Scaffolding Developers' Decision-Making Using the Web. In *Proceedings of the 32nd Annual ACM Symposium on User Interface Software and Technology (UIST '19)*. Association for Computing Machinery, New York, NY, USA, 67–80. <https://doi.org/10.1145/3332165.3347908>
- [32] Michael Xieyang Liu, Aniket Kittur, and Brad A. Myers. 2022. Crystalline: Lowering the Cost for Developers to Collect and Organize Information for Decision Making. In *Proceedings of the 2022 CHI Conference on Human Factors in Computing Systems (CHI '22)*. Association for Computing Machinery, New York, NY, USA, 1–16. <https://doi.org/10.1145/3491102.3501968>
- [33] Michael Xieyang Liu, Tongshuang Wu, Tianying Chen, Franklin Mingzhe Li, Aniket Kittur, and Brad A. Myers. 2024. Selenite: Scaffolding Online Sensemaking with Comprehensive Overviews Elicited from Large Language Models. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*. ACM, Honolulu HI USA, 1–26. <https://doi.org/10.1145/3613904.3642149>
- [34] Antonio Maffei and Alessandro Angrilli. 2019. Spontaneous blink rate as an index of attention and emotion during film clips viewing. *Physiology & Behavior* 204 (May 2019), 256–263. <https://doi.org/10.1016/j.physbeh.2019.02.037>
- [35] Stephann Makri and Lily Buckley. 2020. Down the rabbit hole: Investigating disruption of the information encountering process. *Journal of the Association for Information Science and Technology* 71, 2 (2020), 127–142. <https://doi.org/10.1002/asi.24233> eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1002/asi.24233>
- [36] JAMES G. MILLER. 1960. Information input overload and psychopathology. *American Journal of Psychiatry* 116, 8 (Feb. 1960), 695–704. <https://doi.org/10.1176/ajp.116.8.695> Publisher: American Psychiatric Publishing.
- [37] Muhammad Arief Nugroho, Maman Abdurrohman, Bayu Erfianto, and Mahmud Dwi Sulistiyo. 2024. Meticulous Acquisition System for Tracking User's Natural Kinetics (MAS TUNK): An Approach in Eye Tracking Dataset Collection for Neural Network Training. In *Proceedings of the 2024 Symposium on Eye Tracking Research and Applications (ETRA '24)*. Association for Computing Machinery, New York, NY, USA, 1–8. <https://doi.org/10.1145/3290607.3312801>
- [38] Sharoda A. Paul and Meredith Ringel Morris. 2009. CoSense: enhancing sensemaking for collaborative web search. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '09)*. Association for Computing Machinery, New York, NY, USA, 1771–1780. <https://doi.org/10.1145/1518701.1518974>
- [39] Cole S. Peterson, Jonathan A. Saddler, Natalie M. Halavick, and Bonita Sharif. 2019. A Gaze-Based Exploratory Study on the Information Seeking Behavior of Developers on Stack Overflow. In *Extended Abstracts of the 2019 CHI Conference on Human Factors in Computing Systems (CHI EA '19)*. Association for Computing Machinery, New York, NY, USA, 1–6. <https://doi.org/10.1145/3290607.3312801>
- [40] Adam Rule, Aurélien Tabard, Karen Boyd, and Jim Hollan. 2015. Restoring the Context of Interrupted Work with Desktop Thumbnails. In *37th Annual Meeting of the Cognitive Science Society*. Cognitive Science Society, Pasadena, United States, 1–6. <https://hal.science/hal-01213708>
- [41] Dario D. Salvucci and Joseph H. Goldberg. 2000. Identifying fixations and saccades in eye-tracking protocols. In *Proceedings of the 2000 symposium on Eye tracking research & applications (ETRA '00)*. Association for Computing Machinery, New York, NY, USA, 71–78. <https://doi.org/10.1145/355017.355028>
- [42] m.c. schraefel, Yuxiang Zhu, David Modjeska, Daniel Wigdor, and Shengdong Zhao. 2002. Hunter gatherer: interaction support for the creation and management of within-web-page collections. In *Proceedings of the 11th international conference on World Wide Web (WWW '02)*. Association for Computing Machinery, New York, NY, USA, 172–181. <https://doi.org/10.1145/511446.511469>
- [43] Omar Shaikh, Valentino Emil Chai, Michele Gelfand, Diyi Yang, and Michael S. Bernstein. 2024. Rehearsal: Simulating Conflict to Teach Conflict Resolution. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems (CHI '24)*. Association for Computing Machinery, New York, NY, USA, 1–20. <https://doi.org/10.1145/3613904.3642159>
- [44] Omar Shaikh, Kristina Gligorić, Ashna Khetan, Matthias Gerstgrasser, Diyi Yang, and Dan Jurafsky. 2024. Grounding Gaps in Language Model Generations. <https://doi.org/10.48550/arXiv.2311.09144> arXiv:2311.09144 [cs].
- [45] Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askill, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnson, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. 2023. Towards Understanding Sycophancy in Language Models. <https://doi.org/10.48550/arXiv.2310.13548> arXiv:2310.13548 [cs].
- [46] Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed H. Chi, Nathanael Schärli, and Denny Zhou. 2023. Large Language Models Can Be Easily Distracted by Irrelevant Context. In *Proceedings of the 40th International Conference on Machine Learning*. PMLR, 31210–31227. <https://proceedings.mlr.press/v202/shi23a.html> ISSN: 2640-3498.
- [47] Victoria Smy, Marie Cahillane, and Piers MacLean. 2016. Sensemaking and metacognitive prompting in ill-structured problems. *The International Journal of Information and Learning Technology* 33, 3 (June 2016), 186–199. <https://doi.org/10.1108/IJILT-10-2015-0027> Publisher: Emerald Group Publishing Limited.

- [48] John Sweller. 1988. Cognitive load during problem solving: Effects on learning. *Cognitive Science* 12, 2 (April 1988), 257–285. [https://doi.org/10.1016/0364-0213\(88\)90023-7](https://doi.org/10.1016/0364-0213(88)90023-7)
- [49] Nash Unsworth and Matthew K. Robison. 2015. Individual differences in the allocation of attention to items in working memory: Evidence from pupillometry. *Psychonomic Bulletin & Review* 22, 3 (June 2015), 757–765. <https://doi.org/10.3758/s13423-014-0747-6>
- [50] Nerea Urrestilla and David St-Onge. 2021. Measuring Cognitive Load: Heart-rate Variability and Pupillometry Assessment. In *Companion Publication of the 2020 International Conference on Multimodal Interaction (ICMI '20 Companion)*. Association for Computing Machinery, New York, NY, USA, 405–410. <https://doi.org/10.1145/3395035.3425203>
- [51] Lennart Wachowiak, Peter Tisnikar, Gerard Canal, Andrew Coles, Matteo Leonetti, and Oya Celiktutan. 2022. Analysing Eye Gaze Patterns during Confusion and Errors in Human-Agent Collaborations. In *2022 31st IEEE International Conference on Robot and Human Interactive Communication (RO-MAN)*. 224–229. <https://doi.org/10.1109/RO-MAN53752.2022.9900589> ISSN: 1944-9437.
- [52] Xi Wang, Zoya Bylinskii, Monica Castelhano, James Hillis, and Andrew Duchowski. 2021. EMICS'21: Eye Movements as an Interface to Cognitive State. In *Extended Abstracts of the 2021 CHI Conference on Human Factors in Computing Systems (CHI EA '21)*. Association for Computing Machinery, New York, NY, USA, 1–6. <https://doi.org/10.1145/3411763.3441357>
- [53] Yao Wang, Yue Jiang, Zhiming Hu, Constantin Ruhdorfer, Mihai Băce, and Andreas Bulling. 2024. VisRecall+: Analysing and Predicting Visualisation Recallability from Gaze Behaviour. *Proc. ACM Hum.-Comput. Interact.* 8, ETRA (May 2024), 239:1–239:18. <https://doi.org/10.1145/3655613>
- [54] Dale M. Willows. 1974. Reading between the Lines: Selective Attention in Good and Poor Readers. *Child Development* 45, 2 (1974), 408–415. <https://doi.org/10.2307/1127962> Publisher: [Wiley, Society for Research in Child Development].
- [55] Dale M. Willows and G. E. MacKinnon. 1973. Selective reading: Attention to the "unattended" lines. *Canadian Journal of Psychology / Revue canadienne de psychologie* 27, 3 (1973), 292–304. <https://doi.org/10.1037/h0082480> Place: Canada Publisher: University of Toronto Press.
- [56] Ye Xia, Jinkyu Kim, John Canny, Karl Zipser, Teresa Canas-Bajo, and David Whitney. 2020. Periphery-Fovea Multi-Resolution Driving Model Guided by Human Attention. IEEE Computer Society, 1756–1764. <https://doi.org/10.1109/WACV45572.2020.9093524>

Received 20 February 2025